

Incremental Memory Model for MonetDB

Pavlos Katsogridakis

September 28, 2017

Overview

- ▶ MAL level query memory estimator
- ▶ Groups the MAL instructions into categories
- ▶ Uses previous query traces, and base column statistics

MAL example

```
1 select
2   sum(l_extendedprice * l_discount) as revenue
3 from
4   lineitem
5 where
6   l_shipdate >= date '1994-01-01'
7   and l_shipdate < date '1994-01-01'+interval '1' year
8   and l_discount between 0.06 - 0.01 and 0.06 + 0.01
9   and l_quantity < 24;
10
```

MAL example

```
1  X_16:= bind("sys","lineitem","l_quantity")
2  X_33:= bind("sys","lineitem","l_discount")
3  X_40:= bind("sys","lineitem","l_shipdate")
4  C_13:= tid("sys","lineitem")
5  C_51:= select(X_40,C_13,1994-01-01,1995-01-01)
6  C_65:= select(X_33,C_51,5,7)
7  C_68:= thetaselect(X_16,C_65,2400,"<")
8  X_26:= bind("sys","lineitem","l_extendedprice")
9  X_71:= projection(C_68,X_26)
10 X_72:= projection(C_68,X_33)
11 X_78:= *(X_71,X_72)
12 X_80:= sum(X_78)
13
```

What is memory footprint ?

Definition

The maximum memory usage in the query execution timeline

```
1 max_mem = 0
2 curr_mem = 0
3 for i in ilist:
4     max_mem = max(max_mem, curr_mem + i.mem_fprint)
5     curr_mem = curr_mem + i.mem_fprint - i.free_size
```

Goal

Take the MAL instruction plan and give an estimation of the memory footprint

- ▶ Use only traces from previous executions and basic column statistics

MAL Instruction Grouping

- ▶ Arithmetic (+,-,...)
- ▶ Aggregate (avg, sum, ...)
- ▶ Limit (firstn, sample)
- ▶ Grouping (groupdone)
- ▶ Selection (select, thetaselect)
- ▶ Join
- ▶ Set (intersect, merge, diff)

Output Count Prediction

Arithmetic Operators

Output = Input

Aggregate Operators

Output = 1

Limit Operators

Output = $\min(\text{arg size}, N)$

Set Operators

Merge(A,B)

Output = size(A)+size(B)

Intersect(A,B)

Output = min(size(A),size(B))

Diff(A,B)

Output = size(A)

Select Operators

Range Select

- ▶ Keep a dictionary of all the previous selections
- ▶ Find all the instructions having the same column, op
- ▶ Run a kNN to find the 5 nearest based on the range
- ▶ Extrapolate based on the selectivity and arg size
- ▶ $\text{extrap}(\text{testi}, \text{traini}) = \text{traini.count} * \frac{\text{testi.estim_arg_cnt}}{\text{traini.arg_cnt}} * \frac{\text{testi.range}}{\text{traini.range}}$

Point Select

- ▶ Find all the instructions having the same column, op
- ▶ Run a kNN to find the 5 nearest selects
- ▶ Extrapolate based on the arg size


Join Operators

- ▶ Find all the instructions on the same columns
- ▶ Run a kNN to find the 5 nearest based on the arg size
- ▶ Extrapolate based on the arg size


Tpch sf10 Evaluation

- ▶ Produce 200 queries randomizing the selection values
- ▶ Keep the original query as test set
- ▶ $\text{Error} = 100 * (\text{predict_mem} - \text{real_mem}) / \text{real_mem}$

Tpch sf10 Q6




Tpch sf10



Airtraffic benchmark suite

```
1 SELECT SQL_MIN("Origin", "Dest"),  
2        SQL_MAX("Origin", "Dest") AS route,  
3        COUNT(*) FROM ontime  
4 WHERE 2010-09-11<FlightDate AND FlightDate<2011-09-11  
5 GROUP BY route;
```



```
1 SELECT "DayOfWeek", COUNT(*) AS "Flights"  
2 FROM ontime  
3 WHERE "DepDelay" > 15  
4 GROUP BY "DayOfWeek"  
5 ORDER BY "DayOfWeek";
```




Figure: Q04 memory error


```
1 SELECT "Origin" AS ap, COUNT(*) AS cnt_in
2 FROM ontime
3 WHERE "Dest" = 'ORD'
4 GROUP BY "Origin"
```




Figure: Q10 select error

```

1 SELECT *
2 FROM ontime
3 WHERE "DepDelay" > 15 AND "ArrDelay" > 15
4 AND "Month" = 3 AND "DayofMonth" = 24 AND "Year" = 2013

```




Figure: Q09 select error

Future Work

- ▶ Provide min,avg,max policies
- ▶ Consider parallel executions
- ▶ Handle correlated columns on selections
- ▶ Also use histograms